

Exploratory Data Analysis

Christian Birchler

5/21/2020

Load data and packages

```
library(data.table)
library(ggplot2)
library(grid)
library(gridExtra)
library(ggbiplot)

## Loading required package: plyr
## Loading required package: scales
set.seed(123)

##### READ DATA #####
setwd("~/Desktop/bsc-analysis")
dd.orig <- fread("vm6-no-gc-merged.csv")
idflakies <- fread("idflakies-projects.csv")
#####
```

Clean data

A data set with roughly 400'000 entries by 74 variable has probably some invalid (or corrupted) data entries. First of all they must be identified and treated separately.

```
dd <- dd.orig

# convert all columns with numbers as numeric data type
# collect indices of NA entries -> those are corrupted data entries
dd$deadlock.count <- as.numeric(dd$deadlock.count)

## Warning: NAs introduced by coercion
na.indeces <- which(is.na(dd$deadlock.count))

dd$loaded <- as.numeric(dd$loaded)

## Warning: NAs introduced by coercion
na.indeces <- unique(c(na.indeces, which(is.na(dd$loaded)))))

dd$pools.Metaspase.used <- as.numeric(dd$pools.Metaspase.used)
```

```

## Warning: NAs introduced by coercion
na.indeces <- unique( c(na.indeces, which( is.na(dd$pools.Metaspaces.used)))))

dd$`pools.PS-Eden-Space.committed` <- as.numeric(dd$`pools.PS-Eden-Space.committed`)

## Warning: NAs introduced by coercion
na.indeces <- unique( c(na.indeces, which( is.na(dd$`pools.PS-Eden-Space.committed`)))))

dd$`pools.PS-Eden-Space.init` <- as.numeric(dd$`pools.PS-Eden-Space.init`)

## Warning: NAs introduced by coercion
na.indeces <- unique( c(na.indeces, which( is.na(dd$`pools.PS-Eden-Space.init`)))))

dd$`pools.PS-Eden-Space.max` <- as.numeric(dd$`pools.PS-Eden-Space.max`)

## Warning: NAs introduced by coercion
na.indeces <- unique( c(na.indeces, which( is.na(dd$`pools.PS-Eden-Space.max`)))))

dd$`pools.PS-Eden-Space.used` <- as.numeric(dd$`pools.PS-Eden-Space.used`)

## Warning: NAs introduced by coercion
na.indeces <- unique( c(na.indeces, which( is.na(dd$`pools.PS-Eden-Space.used`)))))

dd$`pools.PS-Survivor-Space.committed` <- as.numeric(dd$`pools.PS-Survivor-Space.committed`)

## Warning: NAs introduced by coercion
na.indeces <- unique( c(na.indeces, which( is.na(dd$`pools.PS-Survivor-Space.committed`)))))

dd$terminated.count <- as.numeric(dd$terminated.count)

## Warning: NAs introduced by coercion
na.indeces <- unique( c(na.indeces, which( is.na(dd$terminated.count)))))

# proportion of corrupted data entries -> ~1%
print(length( na.indeces)/nrow(dd))

## [1] 0.009638677

# omit corrupted data entries
dd <- dd[-na.indeces,]

```

Flakiness

```

unique_tests <- unique(dd[, c('TestCase', 'ProjectName', 'CommitHash')])

# count the different outcomes per test case (group by TestCase, ProjectName and CommitHash)
nr_different_outcomes <- dd[, list(`pass/fail` = `pass/fail`, Count = uniqueN(`pass/fail`))
                           , by=c('TestCase', 'ProjectName', 'CommitHash')]

different_outcomes <- nr_different_outcomes[ Count > 1 , ]

```

```

# less than 20 entries -> build failure -> flaky?
nr_measurements_per_test <- dd[ , list(Count = .N)
                                , by=c('TestCase', 'ProjectName', 'CommitHash') ]

# there are only even number of test measurements of a single test case
# --> there can only build failures if there are less than 20 measurements
less_than_20_entries <- nr_measurements_per_test[ Count < 20 , ] # flaky?

# merge (union) different outcomes with tests which have not 20 measurements
different_outcomes_or_less_20_entries <- merge(different_outcomes, less_than_20_entries
                                                , by=c('TestCase', 'ProjectName', 'CommitHash')
                                                , all=TRUE)

#flaky.own <- different_outcomes_or_less_20_entries
flaky.own <- unique( different_outcomes[, c('TestCase', 'ProjectName', 'CommitHash')] )

# percentage of flaky tests in own measurements
flakiness_rate <- nrow(flaky.own)/nrow(unique_tests)

flakiness_rate

## [1] 0.01381114

```

Here is to consider the definition of flakiness (`flaky.own`). In the current setting I consider a test as flaky if there are different outcomes (pass/failure/error) of the test. I do not consider the case where the test code compilation failed for some but not all iterations.

From all the measurements were 1.3% of the tests flaky according to the definition above.

Comparison with idFlakies data set

```

# set merge with idflakies data set (intersection over 'TestCase', 'CommitHash')
flaky.intersect <- merge(flaky.own, idflakies, by.x=c('TestCase', 'CommitHash')
                           , by.y=c('Test Name', 'SHA'), all=FALSE)

nrow(flaky.intersect)

## [1] 2

```

`flaky.intersect` has only two entries. This means that the recorded flaky tests of the idFlakies data set could not be well replicated. Only two of their flaky tests could be replicated. Nevertheless, from the own measurements multiple additional flaky tests could be identified.

Summary of all software projects

Below is a list of all projects where measurements were taken and the number of flaky and non-flaky tests are also listed.

```

# label data set (flaky/not flaky)
flaky.own$isFlaky <- TRUE
dd.labeled <- merge(dd, flaky.own, by=c('TestCase', 'ProjectName', 'CommitHash')
                     , all=TRUE)

```

```

dd.labeled[ is.na(isFlaky) , "isFlaky"] <- FALSE

# unique tests cases (original data set has before/after entries)
labeled_tests <- unique( dd.labeled[, c('TestCase','ProjectName','CommitHash','isFlaky')] )

flaky_tests_per_project <- labeled_tests[ isFlaky == TRUE , list(NrFlaky = .N)
                                         , by=c('ProjectName','CommitHash') ]
non_flaky_tests_per_project <- labeled_tests[ isFlaky == FALSE , list(NrNotFlaky = .N)
                                         , by=c('ProjectName','CommitHash') ]

projects_overview <- merge(flaky_tests_per_project, non_flaky_tests_per_project,
                            by=c('ProjectName','CommitHash'), all=TRUE)

projects_overview[, -'CommitHash']

##                               ProjectName NrFlaky NrNotFlaky
## 1:                      Activiti      2        171
## 2:                 GoodData-CL     NA         21
## 3:          Java-WebSocket     1       144
## 4:                  Struts      64       268
## 5:           WebCollector     NA         9
## 6:             admiral     NA       136
## 7:        aismessages     NA        44
## 8:          aletheia     NA        31
## 9:         as2-lib     NA        47
## 10:         c2mon      3       372
## 11:          cukes     NA        54
## 12:    db-scheduler     NA        44
## 13: delight-nashorn-sandbox     NA        77
## 14:   elastic-job-lite     NA       553
## 15:            esper     NA         4
## 16:      excelastic     NA        11
## 17: fluent-logger-java     NA        16
## 18:        hadoop     NA       671
## 19:        helios      1       205
## 20: hsac-fitnessse-fixtures     NA       250
## 21:          hutool     NA       508
## 22: incubator-dubbo      10      1509
## 23:   java-cas-client     NA       150
## 24:        jfreechart     NA      2176
## 25:   jhipster-registry      4        49
## 26:          luwak     NA       160
## 27:        marine-api     NA       925
## 28: nexus-repository-helm     NA        18
## 29:      nifi-registry     NA       251
## 30:          noxy     NA        29
## 31:   oci-java-sdk     NA       112
## 32:          one     NA        12
## 33:          orbit      1       184
## 34:          oryx      20       372
## 35:          pippo     NA       207
## 36:    querydsl      16      1841

```

```

## 37:           redpipe    NA     61
## 38:           retrofit   NA    605
## 39:           riptide    8    315
## 40:      rxjava2-extras  NA    385
## 41:           sawmill   NA     2
## 42:           sos       NA     4
## 43:           spring-cloud-aws  NA   296
## 44:  spring-cloud-zuul-ratelimit  NA     3
## 45:           spring-data-ebean  NA    48
## 46:           spring-data-envers  NA    10
## 47:           spring-ws     NA   109
## 48:           timely      72    52
## 49:           tyrus       38   491
## 50:           unix4j      NA   428
## 51:  vertx-completable-future  NA    61
## 52:  vertx-rabbitmq-client    1     5
## 53:           whois      NA  2540
## 54:  wikidata-toolkit     NA   786
## 55:           wildfly     16   118
## 56:  wildfly-maven-plugin   NA    18
## 57:           wro4j      7    861
## 58:           yawp       NA    22
##
##             ProjectName NrFlaky NrNotFlaky

```

Projects of interest

According to the table above it might be more interesting to focus only on projects with a fair amount of flaky tests to gain a better understanding. It does not make really sense to investigate measurements with no flaky tests or only a few. Below you see the depicted projects of interest.

```
projects_overview[c(4, 25, 34, 39, 55), -'CommitHash']
```

```

##             ProjectName NrFlaky NrNotFlaky
## 1:           Struts     64    268
## 2: jhipster-registry     4     49
## 3:           oryx     20    372
## 4:           riptide    8    315
## 5:           wildfly    16   118

```

PCA dimensionality reduction

Since the data set is multivariate it makes sence to do a PCA to see if we can position flaky tests in a region where the data has a high variance. For this purpose I treat each project (listed above) separately.

Struts

```

struts <- dd.labeled[ProjectName == 'Struts',]
# struts$deadlock.count <- as.numeric(struts$deadlock.count)
# struts$loaded <- as.numeric(struts$loaded)
# struts$pools.Metaspace.used <- as.numeric(struts$pools.Metaspace.used)

```

```

# struts$`pools.PS-Eden-Space.committed` <- as.numeric(struts$`pools.PS-Eden-Space.committed`)
# struts$`pools.PS-Eden-Space.init` <- as.numeric(struts$`pools.PS-Eden-Space.init`)
# struts$`pools.PS-Eden-Space.max` <- as.numeric(struts$`pools.PS-Eden-Space.max`)
# struts$`pools.PS-Eden-Space.used` <- as.numeric(struts$`pools.PS-Eden-Space.used`)
# struts$`pools.PS-Survivor-Space.committed` <- as.numeric(struts$`pools.PS-Survivor-Space.committed`)
# struts$terminated.count <- as.numeric(struts$terminated.count)

# waiting count data might be corrupted
#struts$waiting.count <- as.numeric(struts$waiting.count)

#struts[,-c('V1','pass/fail','iteration','TestCase','CommitHash','before/after','ProjectName','deadlock')
ignore <- c('V1','pass/fail','iteration','TestCase','CommitHash','before/after','ProjectName','deadlock'
select_columns <- c()

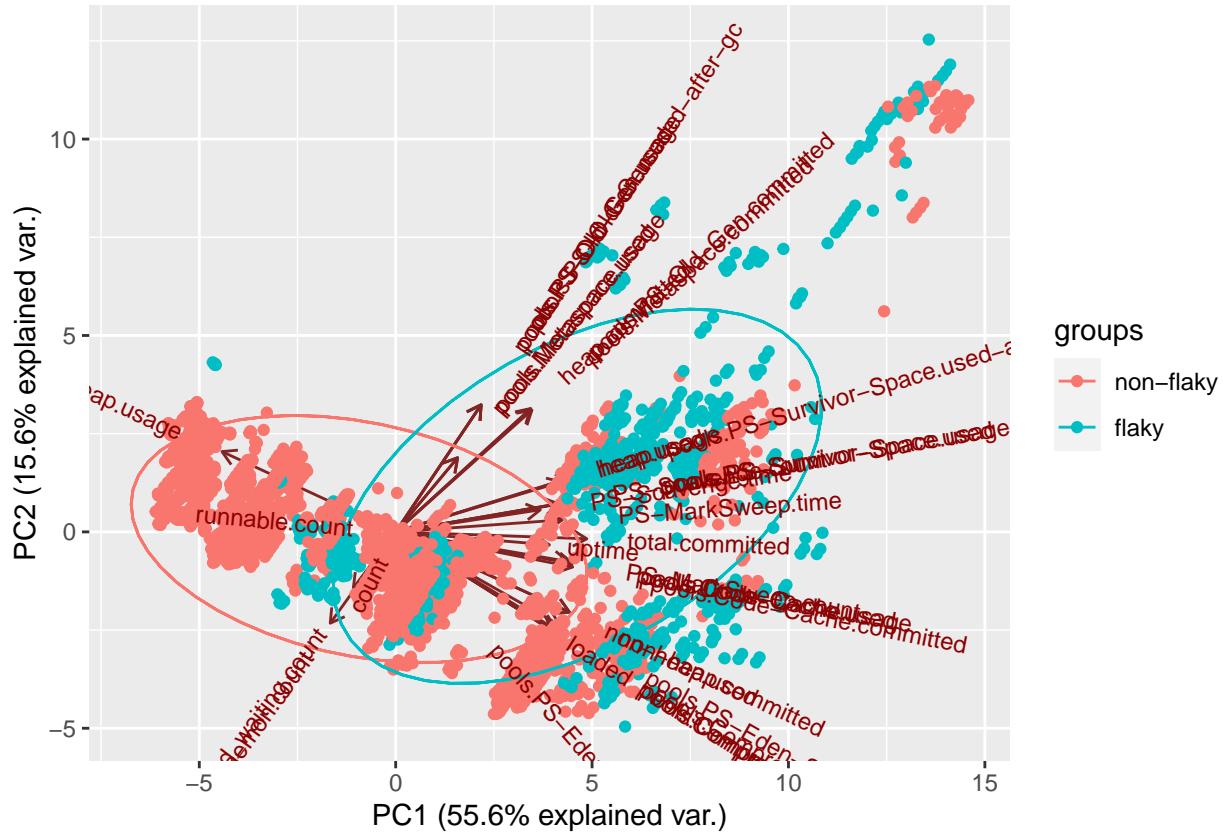
# for-loop cannot handle data.table :-
struts <- as.data.frame(struts)
# PCA cannot be performed if there are vectors if uniform entries
for (i in colnames(struts)){
  x <- struts[3,i]
  if(all(struts[,i] == x) || i %in% ignore){
    next
  }
  select_columns <- c(select_columns, i)
}

#struts <- as.data.table(struts)
#struts[,select_columns]

# perform PCA with scaling and center (avoid missleading eigenvectors)
struts.pca <- prcomp(struts[,select_columns], center=TRUE, scale.=TRUE )

groups <- as.factor(struts$isFlaky)
levels(groups) <- c("non-flaky","flaky")
ggbiplot(struts.pca, groups=groups, ellipse=TRUE, obs.scale=TRUE, var.scale=TRUE)

```



This image shows that the identified flaky tests of ‘Struts’ can be located separately (more or less) within this biplot.

Wildfly

```
wildfly <- dd.labeled[ProjectName == 'wildfly',]
# wildfly$deadlock.count <- as.numeric(wildfly$deadlock.count)
# wildfly$loaded <- as.numeric(wildfly$loaded)
# wildfly$pools.Metaspace.used <- as.numeric(wildfly$pools.Metaspace.used)
# wildfly$pools.PS-Eden-Space.committed` <- as.numeric(wildfly$pools.PS-Eden-Space.committed`)
# wildfly$pools.PS-Eden-Space.init` <- as.numeric(wildfly$pools.PS-Eden-Space.init`)
# wildfly$pools.PS-Eden-Space.max` <- as.numeric(wildfly$pools.PS-Eden-Space.max`)
# wildfly$pools.PS-Eden-Space.used` <- as.numeric(wildfly$pools.PS-Eden-Space.used`)
# wildfly$pools.PS-Survivor-Space.committed` <- as.numeric(wildfly$pools.PS-Survivor-Space.committed`)
# wildfly$terminated.count <- as.numeric(wildfly$terminated.count)

# waiting count data might be corrupted
#struts$waiting.count <- as.numeric(struts$waiting.count)

#struts[,-c('V1','pass/fail','iteration','TestCase','CommitHash','before/after','ProjectName','deadlock')]
#ignore <- c('V1','pass/fail','iteration','TestCase','CommitHash','before/after','ProjectName','deadlock')
select_columns <- c()

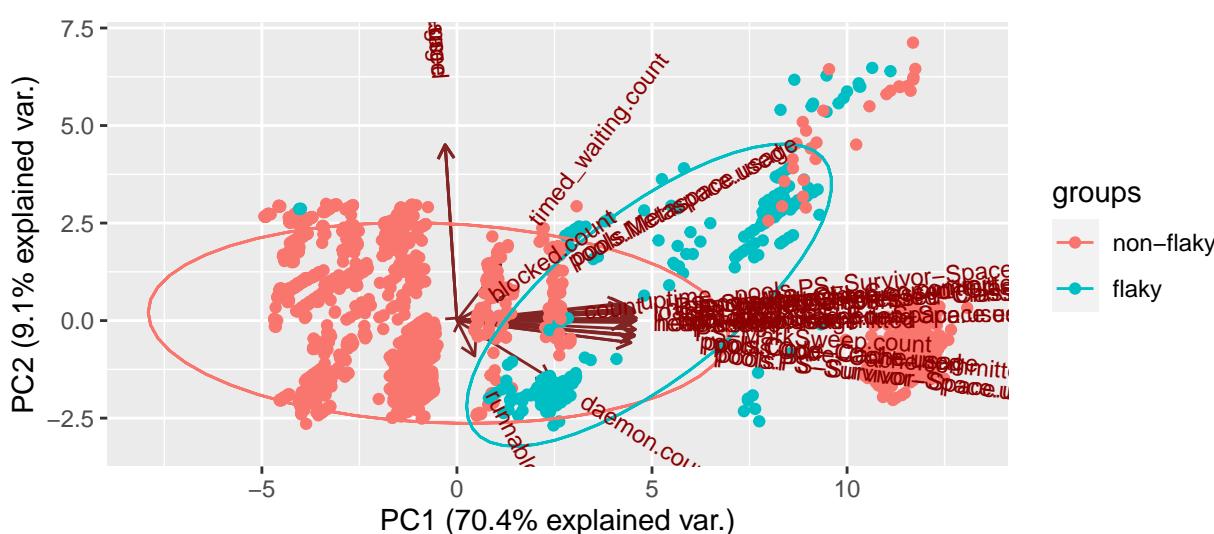
# for-loop cannot handle data.table :-(
```

```
wildfly <- as.data.frame(wildfly)
# PCA cannot be performed if there are vectors if uniform entries
for (i in colnames(wildfly)){
  x <- wildfly[3,i]
  if(all(wildfly[,i] == x) || i %in% ignore){
    next
  }
  select_columns <- c(select_columns, i)
}
}

#struts <- as.data.table(struts)
#struts[,select_columns]

# perform PCA with scaling and center (avoid missleading eigenvectors)
wildfly.pca <- prcomp(wildfly[,select_columns], center=TRUE, scale.=TRUE)

groups <- as.factor(wildfly$isFlaky)
levels(groups) <- c("non-flaky", "flaky")
```



Oryx

```

oryx <- dd.labeled[ProjectName == 'oryx',]
# oryx$deadlock.count <- as.numeric(oryx$deadlock.count)
# oryx$loaded <- as.numeric(oryx$loaded)
# oryx$pools.Metaspace.used <- as.numeric(oryx$pools.Metaspace.used)
# oryx$`pools.PS-Eden-Space.committed` <- as.numeric(oryx$`pools.PS-Eden-Space.committed`)
# oryx$`pools.PS-Eden-Space.init` <- as.numeric(oryx$`pools.PS-Eden-Space.init`)
# oryx$`pools.PS-Eden-Space.max` <- as.numeric(oryx$`pools.PS-Eden-Space.max`)
# oryx$`pools.PS-Eden-Space.used` <- as.numeric(oryx$`pools.PS-Eden-Space.used`)
# oryx$`pools.PS-Survivor-Space.committed` <- as.numeric(oryx$`pools.PS-Survivor-Space.committed`)
# oryx$terminated.count <- as.numeric(oryx$terminated.count)

# waiting count data might be corrupted
struts$waiting.count <- as.numeric(struts$waiting.count)

```

```

#struts[,-c('V1','pass/fail','iteration','TestCase','CommitHash','before/after','ProjectName','deadlock

#ignore <- c('V1','pass/fail','iteration','TestCase','CommitHash','before/after','ProjectName','deadlock
select_columns <- c()

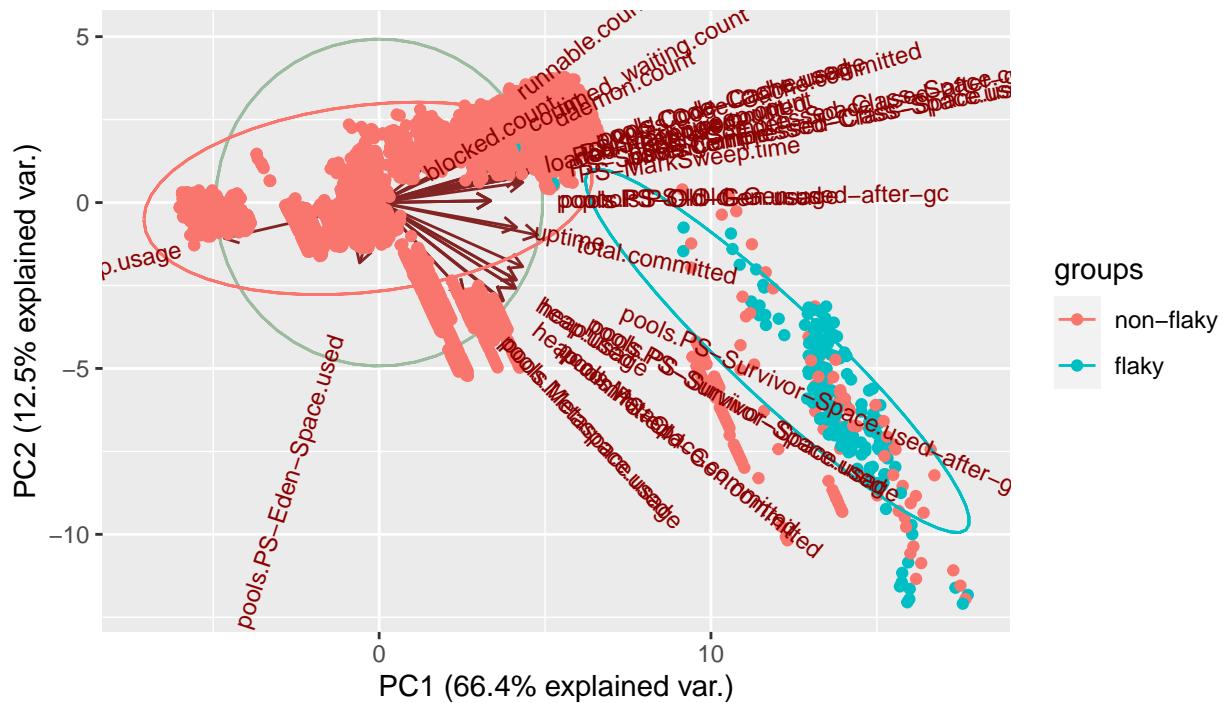
# for-loop cannot handle data.table :-
oryx <- as.data.frame(oryx)
# PCA cannot be performed if there are vectors if uniform entries
for (i in colnames(oryx)){
  x <- oryx[3,i]
  if(all(oryx[,i] == x) || i %in% ignore){
    next
  }
  select_columns <- c(select_columns, i)
}

#struts <- as.data.table(struts)
#struts[,select_columns]

# perform PCA with scaling and center (avoid misleading eigenvectors)
oryx.pca <- prcomp(oryx[,select_columns], center=TRUE, scale.=TRUE )

groups <- as.factor(oryx$isFlaky)
levels(groups) <- c("non-flaky","flaky")
ggbiplots(oryx.pca, groups=groups, ellipse=TRUE, obs.scale=TRUE, var.scale=TRUE, circle=TRUE)

```



jhipster-registry

```
jhipsterregistry <- dd.labeled[ProjectName == 'jhipster-registry',]
# jhipsterregistry$deadlock.count <- as.numeric(jhipsterregistry$deadlock.count)
# jhipsterregistry$loaded <- as.numeric(jhipsterregistry$loaded)
# jhipsterregistry$pools.Metaspace.used <- as.numeric(jhipsterregistry$pools.Metaspace.used)
# jhipsterregistry$pools.PS-Eden-Space.committed` <- as.numeric(jhipsterregistry$pools.PS-Eden-Space.
# jhipsterregistry$pools.PS-Eden-Space.init` <- as.numeric(jhipsterregistry$pools.PS-Eden-Space.init`)
# jhipsterregistry$pools.PS-Eden-Space.max` <- as.numeric(jhipsterregistry$pools.PS-Eden-Space.max`)
# jhipsterregistry$pools.PS-Eden-Space.used` <- as.numeric(jhipsterregistry$pools.PS-Eden-Space.used`)
# jhipsterregistry$pools.PS-Survivor-Space.committed` <- as.numeric(jhipsterregistry$pools.PS-Survivor
# jhipsterregistry$terminated.count <- as.numeric(jhipsterregistry$terminated.count)

# waiting count data might be corrupted
#struts$waiting.count <- as.numeric(struts$waiting.count)

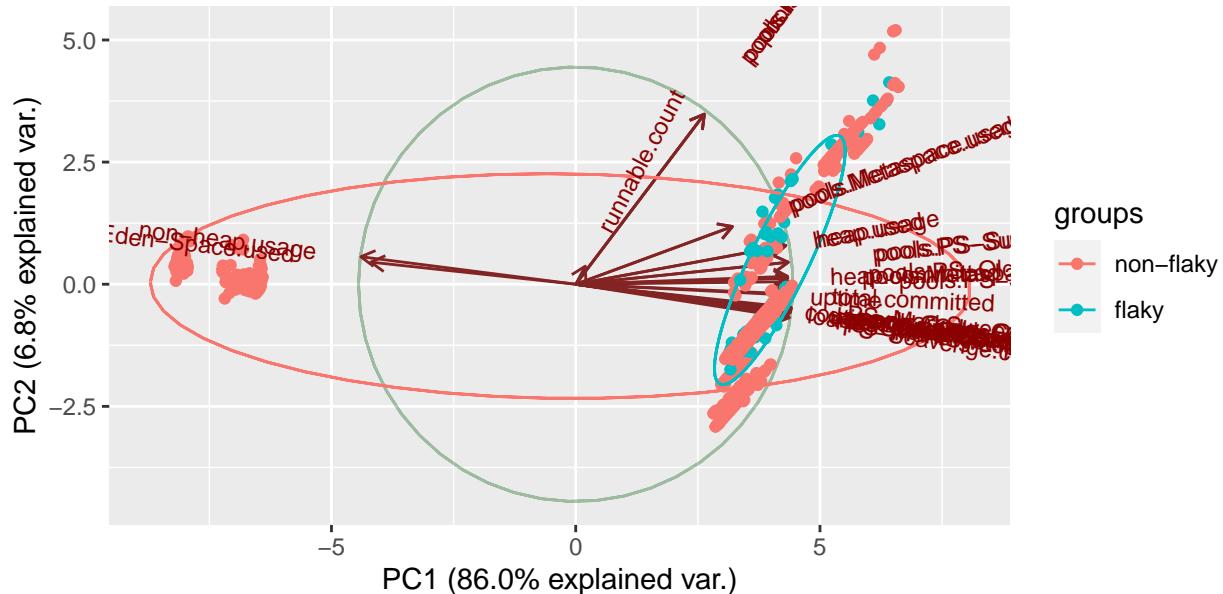
#struts[,-c('V1','pass/fail','iteration','TestCase','CommitHash','before/after','ProjectName','deadlock
#ignore <- c('V1','pass/fail','iteration','TestCase','CommitHash','before/after','ProjectName','deadloc
select_columns <- c()

# for-loop cannot handle data.table :-( 
jhipsterregistry <- as.data.frame(jhipsterregistry)
# PCA cannot be performed if there are vectors if uniform entries
for (i in colnames(jhipsterregistry)){
  x <- jhipsterregistry[3,i]
  if(all(jhipsterregistry[,i] == x) || i %in% ignore){
    next
  }
  select_columns <- c(select_columns, i)
}

#struts <- as.data.table(struts)
#struts[,select_columns]

# perform PCA with scaling and center (avoid misleading eigenvectors)
jhipsterregistry.pca <- prcomp(jhipsterregistry[,select_columns], center=TRUE, scale.=TRUE )

groups <- as.factor(jhipsterregistry$isFlaky)
levels(groups) <- c("non-flaky","flaky")
ggbiplots(jhipsterregistry.pca, groups=groups, ellipse=TRUE, obs.scale=TRUE, var.scale=TRUE, circle=TRUE)
```



This biplot does not show a clear separate cluster of flaky tests. The flaky tests are evenly distributed among non flaky tests on the right side. It might be that for this project too few data are available in comparison to the previous three projects and their PCA biplot.

PCA of all all measurements together

Now we take a look how the biplots look like if we do not separate the projects. For this purpose it makes sense also to scale (as usual) the data for the PCA.

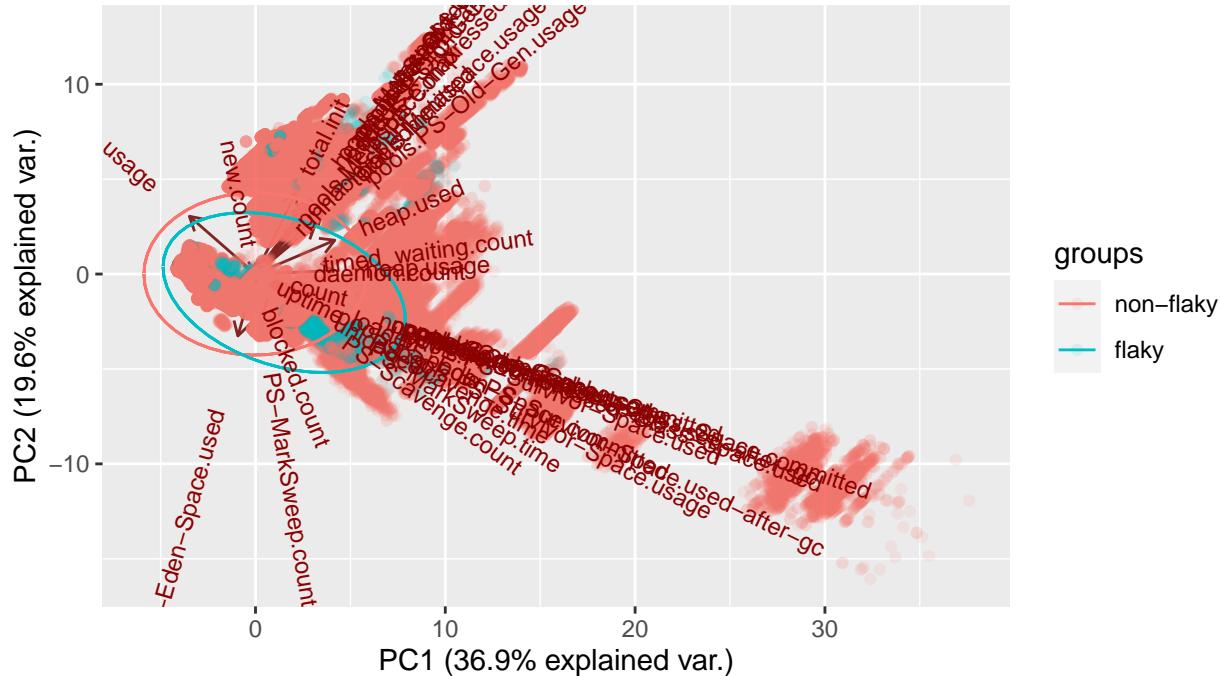
```
# convert all columns with numbers as numeric data type
# collect indices of NA entries -> those are corrupted data entries
all_together <- dd.labeled

select_columns <- c()

# for-loop cannot handle data.table :-( 
all_together <- as.data.frame(all_together)
# PCA cannot be performed if there are vectors of uniform entries
for (i in colnames(all_together)){
  x <- all_together[3,i]
  if(all(all_together[,i] == x) || i %in% ignore){
    next
  }
  select_columns <- c(select_columns, i)
}

# perform PCA with scaling and center (avoid misleading eigenvectors)
all_together.pca <- prcomp(all_together[,select_columns], center=TRUE, scale.=TRUE)

groups <- as.factor(all_together$isFlaky)
levels(groups) <- c("non-flaky", "flaky")
ggbiplots(all_together.pca, groups=groups, ellipse=TRUE, obs.scale=TRUE, var.scale=TRUE, alpha=0.1)
```



This plot with all data points is pretty useless since there are too many data points which overlap heavily! A way to tackle is to downsample the data set. Try to balance the data. Note that only 1.4% of the data are flaky.

```
nr_non_flaky <- 120
nr_flaky <- 40

non_flaky_sample <- dd.labeled[sample( which( !dd.labeled$isFlaky), nr_non_flaky), ]
flaky_sample <- dd.labeled[sample( which( dd.labeled$isFlaky), nr_flaky), ]

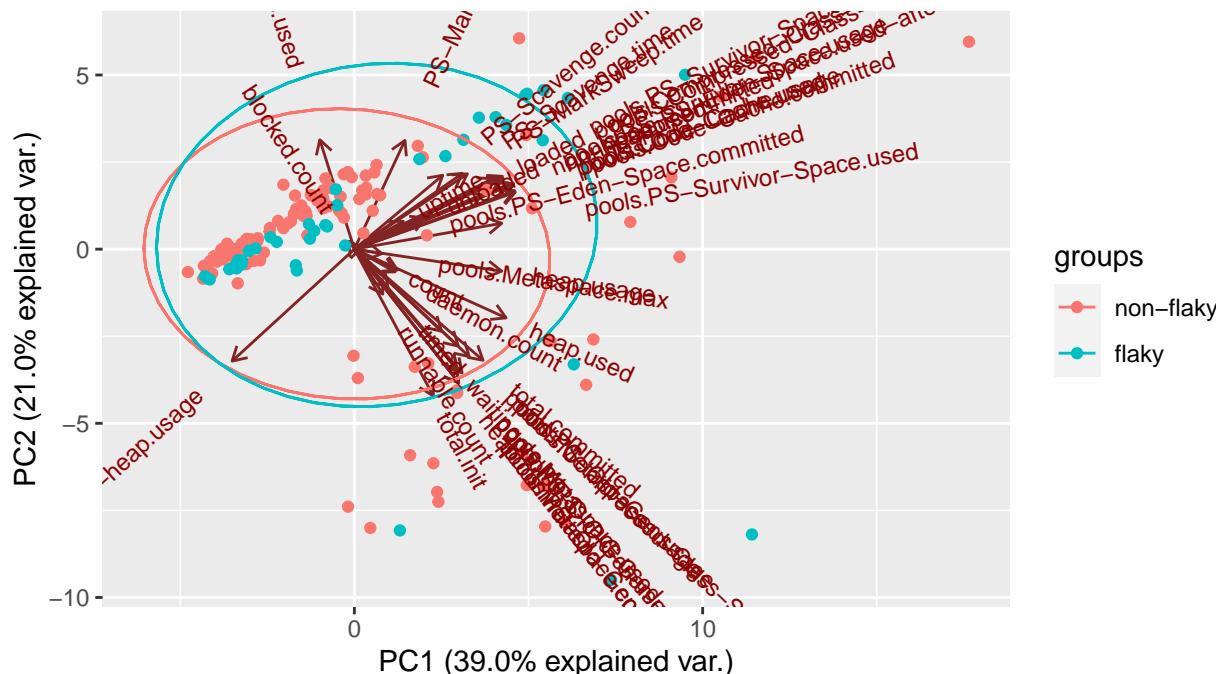
all_together.downsample <- rbind(non_flaky_sample, flaky_sample)

select_columns <- c()

# for-loop cannot handle data.table :(
all_together.downsample <- as.data.frame(all_together.downsample)
# PCA cannot be performed if there are vectors of uniform entries
for (i in colnames(all_together.downsample)){
  x <- all_together.downsample[3,i]
  if(all(all_together.downsample[,i] == x) || i %in% ignore){
    next
  }
  select_columns <- c(select_columns, i)
}

# perform PCA with scaling and center (avoid misleading eigenvectors)
all_together.downsample.pca <- prcomp(all_together.downsample[,select_columns], center=TRUE, scale.=TRUE)

groups <- as.factor(all_together.downsample$isFlaky)
levels(groups) <- c("non-flaky","flaky")
ggbiplots(all_together.downsample.pca, groups=groups, ellipse=TRUE, obs.scale=TRUE, var.scale=TRUE)
```



A unique cluster of flaky tests cannot be identified if we consider all test of all projects together. I think one of the main insights of all the biplots so far is that flaky tests of a project build a cluster (in biplots) but only within a certain project. In the case where arbitrarily tests of different projects were chosen it is very unlikely to cluster them.

Reported flaky tests of the idFlaky data set

```

dd.idflakies <- merge( dd, idflakies, by.x=c('TestCase','CommitHash'), by.y=c('Test Name','SHA'), all=FALSE)

nr_measure_idflakies_entries <- nrow( unique( dd.idflakies[, c('TestCase','CommitHash')])))
# proportion of measured idlfakies entries
nr_measure_idflakies_entries/nrow(idflakies)

## [1] 0.5379147

# label idflakies data set with isFlaky
dd.idflakies$isFlaky <- TRUE

# merge dd.idflakies with all non flaky tests
dd.labeled <- as.data.frame( dd.labeled)
dd.idflakies <- as.data.frame( dd.idflakies)
tmp <- merge( dd.labeled[ which( dd.labeled$isFlaky == FALSE),], dd.idflakies[, c('TestCase','CommitHash')],
              by=c('TestCase','CommitHash'), all=TRUE, allow.cartesian=FALSE)

tmp <- unique( as.data.table(tmp))

tmp[isFlaky.y == TRUE, 'isFlaky.x'] <- TRUE
tmp$isFlaky.y <- NULL
tmp <- tmp[,1:74]
colnames(tmp)[length(colnames(tmp))] <- 'isFlaky'

```

```

# data set of interest
idflakies_only <- na.omit(tmp)
rm(tmp)

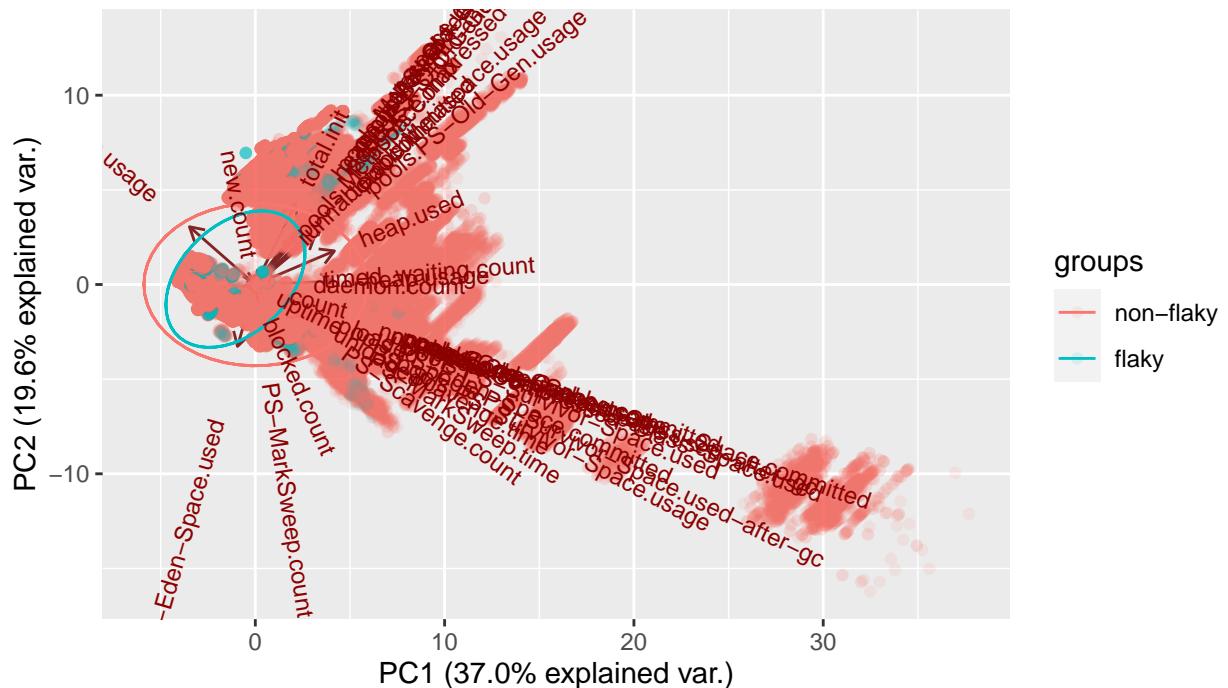
select_columns <- c()

# for-loop cannot handle data.table :-( 
idflakies_only <- as.data.frame(idflakies_only)
# PCA cannot be performed if there are vectors of uniform entries
for (i in colnames(idflakies_only)){
  x <- idflakies_only[3,i]
  if(all(idflakies_only[,i] == x) || i %in% ignore){
    next
  }
  select_columns <- c(select_columns, i)
}

# perform PCA with scaling and center (avoid misleading eigenvectors)
idflakies_only.pca <- prcomp(idflakies_only[,select_columns], center=TRUE, scale.=TRUE )

groups <- as.factor(idflakies_only$isFlaky)
levels(groups) <- c("non-flaky","flaky")
ggbiplots(idflakies_only.pca, groups=groups, ellipse=TRUE, obs.scale=TRUE, var.scale=TRUE, alpha=0.1)

```



Also in this case we can not clearly identify a structure of flaky tests. It is similar like above where all measurements were considered. Let's differentiate again within certain projects.

Within projects with idFlakies flaky tests only

Below you see a list of the available measurements (not tests)

	ProjectName	flakyMeasurements	notflakyMeasurements
## 1:	c2mon	0	7432
## 2:	orbit	20	3700
## 3:	WebCollector	20	160
## 4:	hutool	20	10140
## 5:	oryx	20	7188
## 6:	excelastic	20	200
## 7:	esper	0	80
## 8:	rxjava2-extras	20	7680
## 9:	db-scheduler	40	920
## 10:	GoodData-CL	20	400
## 11:	as2-lib	0	940
## 12:	one	0	240
## 13:	spring-cloud-zuul-ratelimit	0	60
## 14:	oci-java-sdk	0	2240
## 15:	aletheia	60	560
## 16:	querydsl	0	25242
## 17:	noxy	40	540
## 18:	helios	0	4034
## 19:	admiral	0	2830
## 20:	delight-nashorn-sandbox	40	1500
## 21:	aismessages	40	840
## 22:	sos	0	100
## 23:	elastic-job-lite	200	10960
## 24:	jhipster-registry	20	610
## 25:	sawmill	0	40
## 26:	vertx-rabbitmq-client	20	80
## 27:	yawp	20	420
## 28:	cukes	20	1060
## 29:	vertx-completable-future	1000	2180
## 30:	redpipe	100	1120
## 31:	whois	0	50800
## 32:	marine-api	40	18460
## 33:	hsac-fitnessse-fixtures	20	5000
## 34:	Activiti	0	3122
## 35:	incubator-dubbo	440	35824
## 36:	hadoop	128	3392
## 37:	nifi-registry	0	5020
## 38:	Struts	48	4226
## 39:	fluent-logger-java	20	300
## 40:	tyrus	20	10140
## 41:	java-cas-client	20	3000
## 42:	Java-WebSocket	1020	1860
## 43:	wildfly	860	1322
## 44:	jfreechart	20	43500

```
## 45:      nexusrepository-helm          20           340
## 46:      spring-cloud-aws             20          5900
## 47:      spring-data-ebean            20           940
## 48:      spring-data-envers           20           180
## 49:      spring-ws                   0            2180
## 50:      unix4j                      20          8540
## 51:      wikidata-toolkit            100          15620
## 52:      wildfly-maven-plugin         0            306
## 53:      riptide                     40          28974
## 54:      retrofit                    180          11920
## 55:      wro4j                       20          17424
## 56:      pippo                      100          4220
## 57:      timely                      6            504
## 58:      luwak                      0            3360
##
##                                         ProjectName flakyMeasurements notflakyMeasurements
```